

Machine Learning Models applied to IoT Threat Classification

1. Introduction

The Internet of Things (IoT) has made everyday life easier by connecting smart devices, improving efficiency, and automating tasks. However, as more devices connect to the internet, security threats also increase. Hackers can exploit vulnerabilities in IoT networks, leading to data breaches, system failures, and cyberattacks. To keep IoT systems secure, we need a reliable threat detection system that can identify and prevent attacks.

In this project, we use machine learning to detect IoT threats by analyzing network traffic data. First, we load and preprocess the dataset, ensuring that the data is clean and well-structured. We then perform feature selection using a method called SelectKBest, which helps us choose the most important factors in identifying threats.

Next, we train and evaluate different machine learning models, including Logistic Regression, Decision Trees, and Neural Networks. Each model is tested for accuracy, precision, and recall to determine its effectiveness. To further improve our detection system, we create an ensemble model, which combines the best-performing models to enhance accuracy and reliability.

By the end of this project, we demonstrate that machine learning can effectively detect IoT threats, and an ensemble approach provides even better results. This study highlights the importance of feature selection, model evaluation, and ensemble techniques in improving security in IoT networks. Future research can explore deep learning methods for even more advanced threat detection.

2. Base Libraries

In this section, the necessary libraries and modules are imported. These range from data manipulation libraries like pandas and numpy to visualization tools such as matplotlib and seaborn. Additionally, various machine learning models and evaluation metrics from the Scikit-learn library are also imported. To ensure a smooth execution, any potential warnings, except for deprecation warnings, are suppressed.

```
# Useful Libraries
import pandas as pd
import numpy as np
import warnings
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

```

# Feature selection
from sklearn.feature_selection import SelectKBest, f_classif

# Metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score

# Statistic and machine learning models
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier, AdaBoostClassifier, StackingClassifier,
ExtraTreesClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neural_network import MLPClassifier

# Ignore all warnings
warnings.filterwarnings("ignore")

# Re-enable DeprecationWarning
warnings.filterwarnings("default", category=DeprecationWarning)

```

3. Data Loading and Preprocessing

The dataset, sourced from a CSV file, is loaded into a DataFrame for inspection and manipulation. Preliminary exploration includes examining the first few rows, understanding the data's structure, and visualizing the distribution of labels. Any missing values are handled to ensure data integrity.

```

# Load the CSV file into a DataFrame
df = pd.read_csv('data/DDoS_part-00000-363d1ba3-8ab5-4f96-bc25-4d5862db7cb9-
c000.csv', sep=',')

# Display the first 10 rows of the DataFrame
df.head(10)

```

	flow_duration	Header_Length	Protocol	Type	Duration	Rate	\
0	0.000000	54.00		6.00	64.00	0.329807	
1	0.000000	0.00		1.00	64.00	33.396799	
2	1.052463	108.00		6.00	64.00	1.902353	
3	0.000000	54.20		6.00	64.00	11.243547	
4	0.223192	61.54		6.11	64.64	9.087882	
5	0.000000	54.00		6.00	64.00	17.333181	

6	0.000000	0.00	1.00	75.46	0.000000
7	0.000000	0.00	1.00	64.00	1.507148
8	0.088335	29216.00	17.00	64.00	7752.316374
9	0.000000	54.00	6.00	64.00	50.099188

	Srate	Drate	fin_flag_number	syn_flag_number	rst_flag_number	...
\						
0	0.329807	0.0	1.0	0.0	1.0	...
1	33.396799	0.0	0.0	0.0	0.0	...
2	1.902353	0.0	0.0	1.0	0.0	...
3	11.243547	0.0	0.0	1.0	0.0	...
4	9.087882	0.0	0.0	0.0	0.0	...
5	17.333181	0.0	0.0	0.0	0.0	...
6	0.000000	0.0	0.0	0.0	0.0	...
7	1.507148	0.0	0.0	0.0	0.0	...
8	7752.316374	0.0	0.0	0.0	0.0	...
9	50.099188	0.0	0.0	1.0	0.0	...

	Std	Tot size	IAT	Number	Magnitue	Radius	Covariance
\							
0	0.000000	54.00	8.334383e+07	9.5	10.392305	0.000000	0.000000
1	0.000000	42.00	8.312799e+07	9.5	9.165151	0.000000	0.000000
2	0.000000	54.00	8.336548e+07	9.5	10.392305	0.000000	0.000000
3	0.619849	54.20	8.308906e+07	9.5	10.409168	0.878113	3.254011
4	1.692073	54.77	8.333087e+07	9.5	10.434347	2.398780	32.140680
5	0.000000	54.00	8.307592e+07	9.5	10.392305	0.000000	0.000000
6	0.000000	42.00	8.315005e+07	9.5	9.165151	0.000000	0.000000
7	0.000000	42.00	8.315032e+07	9.5	9.165151	0.000000	0.000000
8	0.000000	50.00	8.310233e+07	9.5	10.000000	0.000000	0.000000
9	0.000000	54.00	8.309398e+07	9.5	10.392305	0.000000	0.000000

	Variance	Weight	label
0	0.00	141.55	DDoS-RSTFINFlood
1	0.00	141.55	DDoS-ICMP_Flood
2	0.00	141.55	DDoS-SynonymousIP_Flood
3	0.12	141.55	DDoS-SYN_Flood
4	0.09	141.55	DDoS-PSHACK_Flood
5	0.00	141.55	DDoS-TCP_Flood
6	0.00	141.55	DDoS-ICMP_Flood
7	0.00	141.55	DDoS-ICMP_Flood
8	0.00	141.55	DDoS-UDP_Flood
9	0.00	141.55	DDoS-SYN_Flood

[10 rows x 47 columns]

```
# Display a concise summary of the DataFrame's columns, non-null values, and
data types
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 173777 entries, 0 to 173776  
Data columns (total 47 columns):
```

#	Column	Non-Null Count	Dtype
0	flow_duration	173777 non-null	float64
1	Header_Length	173777 non-null	float64
2	Protocol Type	173777 non-null	float64
3	Duration	173777 non-null	float64
4	Rate	173777 non-null	float64
5	Srate	173777 non-null	float64
6	Drate	173777 non-null	float64
7	fin_flag_number	173777 non-null	float64
8	syn_flag_number	173777 non-null	float64
9	rst_flag_number	173777 non-null	float64
10	psh_flag_number	173777 non-null	float64
11	ack_flag_number	173777 non-null	float64
12	ece_flag_number	173777 non-null	float64
13	cwr_flag_number	173777 non-null	float64
14	ack_count	173777 non-null	float64
15	syn_count	173777 non-null	float64
16	fin_count	173777 non-null	float64
17	urg_count	173777 non-null	float64
18	rst_count	173777 non-null	float64
19	HTTP	173777 non-null	float64
20	HTTPS	173777 non-null	float64
21	DNS	173777 non-null	float64
22	Telnet	173777 non-null	float64
23	SMTP	173777 non-null	float64
24	SSH	173777 non-null	float64
25	IRC	173777 non-null	float64
26	TCP	173777 non-null	float64
27	UDP	173777 non-null	float64
28	DHCP	173777 non-null	float64
29	ARP	173777 non-null	float64
30	ICMP	173777 non-null	float64
31	IPv	173777 non-null	float64
32	LLC	173777 non-null	float64
33	Tot sum	173777 non-null	float64
34	Min	173777 non-null	float64
35	Max	173777 non-null	float64
36	AVG	173777 non-null	float64
37	Std	173777 non-null	float64
38	Tot size	173777 non-null	float64
39	IAT	173777 non-null	float64
40	Number	173777 non-null	float64
41	Magnitue	173777 non-null	float64
42	Radius	173777 non-null	float64
43	Covariance	173777 non-null	float64
44	Variance	173777 non-null	float64

```

45 Weight          173777 non-null float64
46 label           173777 non-null object
dtypes: float64(46), object(1)
memory usage: 62.3+ MB

```

Get the number of rows and columns in the DataFrame

```

df.shape

(173777, 47)

```

Display statistical summary of the DataFrame's columns

```
df.describe()
```

	flow_duration	Header_Length	Protocol Type	Duration \
count	173777.000000	1.737770e+05	173777.000000	173777.000000
mean	0.341774	6.937612e+03	6.723544	64.287490
std	13.149956	4.107415e+04	4.992523	2.582139
min	0.000000	0.000000e+00	0.770000	19.210000
25%	0.000000	5.375000e+01	5.940000	64.000000
50%	0.000000	5.400000e+01	6.000000	64.000000
75%	0.044044	9.612000e+01	6.000000	64.000000
max	3990.334709	3.188817e+06	17.000000	211.070000

	Rate	Srate	Drate	fin_flag_number \
count	1.737770e+05	1.737770e+05	173777.000000	173777.000000
mean	9.493362e+03	9.493362e+03	0.000007	0.11883
std	1.038008e+05	1.038008e+05	0.002069	0.32359
min	0.000000e+00	0.000000e+00	0.000000	0.000000
25%	2.077204e+00	2.077204e+00	0.000000	0.000000
50%	1.356388e+01	1.356388e+01	0.000000	0.000000
75%	7.288358e+01	7.288358e+01	0.000000	0.000000
max	5.242880e+06	5.242880e+06	0.848465	1.00000

	syn_flag_number	rst_flag_number	...	AVG	Std \
count	173777.000000	173777.000000	...	173777.000000	173777.000000
mean	0.223453	0.121443	...	78.098595	18.339989
std	0.416561	0.326642	...	147.318503	95.785662
min	0.000000	0.000000	...	42.000000	0.000000
25%	0.000000	0.000000	...	50.000000	0.000000
50%	0.000000	0.000000	...	54.000000	0.000000
75%	0.000000	0.000000	...	54.000000	0.000000
max	1.000000	1.000000	...	2279.491224	1350.157783

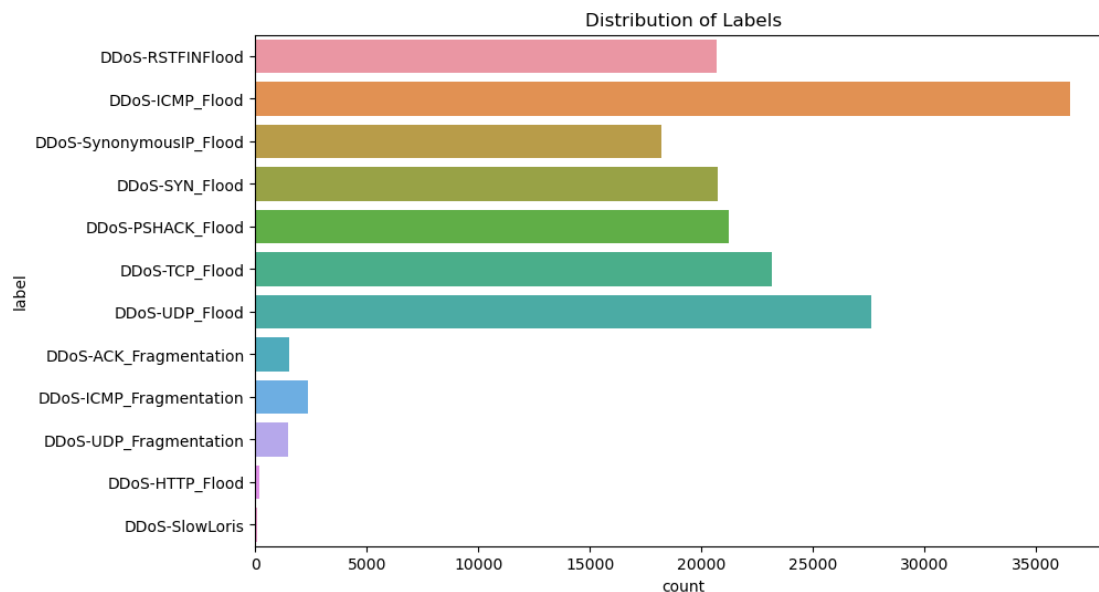
	Tot size	IAT	Number	Magnitue \
count	173777.000000	1.737770e+05	173777.000000	173777.000000
mean	78.138102	8.319127e+07	9.498836	11.126329
std	147.257117	1.340093e+06	0.068870	5.578303

min	42.000000	0.000000e+00	1.000000	9.165151
25%	50.000000	8.309760e+07	9.500000	10.000000
50%	54.000000	8.312886e+07	9.500000	10.392305
75%	54.000000	8.334092e+07	9.500000	10.392305
max	2338.910000	1.001943e+08	10.000000	63.519115

	Radius	Covariance	Variance	Weight
count	173777.000000	1.737770e+05	173777.000000	173777.000000
mean	25.931840	1.079894e+04	0.059744	141.521179
std	135.451819	6.465186e+04	0.177614	1.464649
min	0.000000	0.000000e+00	0.000000	1.000000
25%	0.000000	0.000000e+00	0.000000	141.550000
50%	0.000000	0.000000e+00	0.000000	141.550000
75%	0.000000	0.000000e+00	0.000000	141.550000
max	1912.134518	9.226184e+06	0.950000	141.550000

[8 rows x 46 columns]

```
# Plotting the distribution of the 'label' column
plt.figure(figsize=(10, 6))
sns.countplot(y=df['label'])
plt.title('Distribution of Labels')
plt.show()
```



```
# Check for null values
print(df.isnull().sum())
```

Drop rows with null values (you can also fill them if you prefer)
df.dropna(inplace=True)

flow_duration	0
Header_Length	0
Protocol Type	0
Duration	0
Rate	0
Srate	0
Drate	0
fin_flag_number	0
syn_flag_number	0
rst_flag_number	0
psh_flag_number	0
ack_flag_number	0
ece_flag_number	0
cwr_flag_number	0
ack_count	0
syn_count	0
fin_count	0
urg_count	0
rst_count	0
HTTP	0
HTTPS	0
DNS	0
Telnet	0
SMTP	0
SSH	0
IRC	0
TCP	0
UDP	0
DHCP	0
ARP	0
ICMP	0
IPv	0
LLC	0
Tot sum	0
Min	0
Max	0
AVG	0
Std	0
Tot size	0
IAT	0
Number	0
Magnitue	0
Radius	0
Covariance	0
Variance	0
Weight	0

```
label          0
dtype: int64
```

4. Feature Selection

To improve the efficiency and accuracy of our models, we employ a feature selection method to determine and retain only the most relevant features. This process is essential to reduce the dimensionality of the dataset and improve model training times.

```
# Create a new DataFrame 'X' by dropping the 'label' column
X = df.drop('label', axis=1)

# Assign the 'label' column values to the variable 'y'
y = df['label']

# Initialize a SelectKBest object to select the top 10 features based on the ANOVA F-value
selector = SelectKBest(f_classif, k=10)

# Fit the selector to the data and transform 'X' to retain only the top 10 features
X_new = selector.fit_transform(X, y)

# Get the column names of the selected features from 'X'
selected_features = X.columns[selector.get_support()]

# Print the names of the selected features
print("Selected Features:", selected_features)

Selected Features: Index(['Protocol Type', 'fin_flag_number',
                          'syn_flag_number',
                          'rst_flag_number', 'psh_flag_number', 'ack_count', 'fin_count', 'TCP',
                          'UDP', 'ICMP'],
                        dtype='object')

# Create a new DataFrame with the selected features and the label
df = pd.concat([pd.DataFrame(X_new, columns=selected_features),
                y.reset_index(drop=True)], axis=1)
```


5. Model Training and Validation

Here, we dive into the heart of our analysis. Each machine learning model is trained on a subset of the data and then validated on a separate set. By comparing the models' performances based on accuracy, precision, and recall, we identify the top-performing models. These models are then combined into an ensemble model, aiming to leverage their combined strengths for optimal threat detection.

```
# Create a new DataFrame 'X' by dropping the 'label' column
X = df.drop('label', axis=1)

# Assign the 'label' column values to the variable 'y'
y = df['label']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

This code block sets up, trains, and evaluates a variety of machine learning models on a given dataset. Initially, a dictionary named "models" is defined, where each key is the name of a machine learning algorithm and its associated value is an instance of that algorithm's classifier. Algorithms such as Logistic Regression, Decision Trees, Random Forest, and KNN, among others, are included. The classifiers are initialized with specific parameters, like a maximum number of iterations for the Logistic Regression and Neural Network models.

In the subsequent portion of the code, each model from the "models" dictionary is trained on a training dataset, X_train and y_train. After training, predictions are made on a test dataset, X_test. The accuracy, precision, and recall of these predictions, relative to the true values y_test, are then computed. The results for each model, including its name and evaluation metrics, are stored in the "results" list. This approach provides a straightforward way to compare the performance of different machine learning models on the same dataset.

```
# Define the models
models = {
    'Logistic Regression': LogisticRegression(max_iter=500),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(),
    'Gradient Boosting': GradientBoostingClassifier(),
    'K-Nearest Neighbors': KNeighborsClassifier(),
}

results = []

# Train, test and evaluate each model
for name, model in models.items():
```

```

model.fit(X_train, y_train)
y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')

results.append({
    'Model': name,
    'Accuracy': accuracy,
    'Precision': precision,
    'Recall': recall
})

```

The part performs model ensembling, a technique where multiple machine learning models are combined to achieve better predictive performance than any individual model. First, it evaluates the performance of various standalone models on a dataset and identifies the top three models based on accuracy.

Once these top models are determined, a stacking ensemble is constructed. In stacking, predictions from individual models are used as inputs to another model (the meta-model), which then makes the final predictions. After training this ensemble model on the training data, it is evaluated on the test data, and its performance metrics (accuracy, precision, and recall) are computed. The aim is to capitalize on the strengths of the top models and possibly achieve superior results through their combined predictions.

```

# Convert the results to a DataFrame and sort by Accuracy
results_df = pd.DataFrame(results).sort_values(by='Accuracy',
ascending=False)

# Get the top 3 models
top_models = results_df.head(3)['Model'].tolist()

# Print the names of the models used in the ensemble
print(f"Models used in the ensemble: {'', '.join(top_models)}")

# Create the ensemble model using the top 3 models
ensemble_model = StackingClassifier(estimators=[(name, models[name]) for name
in top_models])
ensemble_model.fit(X_train, y_train)
y_pred_ensemble = ensemble_model.predict(X_test)

# Evaluate the ensemble model
accuracy_ensemble = accuracy_score(y_test, y_pred_ensemble)
precision_ensemble = precision_score(y_test, y_pred_ensemble,
average='weighted')

```

```
recall_ensemble = recall_score(y_test, y_pred_ensemble, average='weighted')
```

```
# Add the ensemble model's results to the results list
```

```
results.append({  
    'Model': 'Ensemble Model',  
    'Accuracy': accuracy_ensemble,  
    'Precision': precision_ensemble,  
    'Recall': recall_ensemble  
})
```

Models used in the ensemble: K-Nearest Neighbors, Neural Network, Gradient Boosting

```
# Convert the results to a DataFrame and sort by Accuracy
```

```
final_results_df = pd.DataFrame(results).sort_values(by='Accuracy',  
ascending=False)
```

```
# Display the results
```

```
final_results_df.head(10)
```

	Model	Accuracy	Precision	Recall
5	Ensemble Model	0.877230	0.873828	0.877230
2	Random Forest	0.877028	0.856051	0.877028
3	Gradient Boosting	0.876913	0.927239	0.876913
1	Decision Tree	0.876511	0.858782	0.876511
0	Logistic Regression	0.874727	0.817484	0.874727
4	K-Nearest Neighbors	0.872799	0.860573	0.872799

Ensemble Model: This model tops the list with an accuracy of approximately 87.72%. The ensemble technique, which combines predictions from multiple models to improve overall performance, seems to have paid off. Its precision score is slightly higher than its accuracy, suggesting that the model's predictions are quite reliable.

Random Forest: Random Forest comes in a close second, with an accuracy almost matching the ensemble model. Its precision is notably less than its accuracy, which means that among the instances it predicted positively, a high percentage were incorrect.

Gradient Boosting: This model's precision stands out, being the highest among the top models. Although its accuracy is slightly lower than the Random Forest, its high precision suggests that its positive predictions are very reliable.

Decision Tree: The Decision Tree model's performance metrics are closely aligned with the Gradient Boosting. This is expected since Random Forests are essentially an ensemble of Decision Trees.

Logistic Regression: This model has similar accuracy scores, but their precision is lower than the models ranked above them. This might imply that these models are making more false positive predictions compared to the top-performing models.

K-Nearest Neighbors (KNN): This model has the lowest accuracy score among the models. It performs well in terms of accuracy, but its precision score suggest there's room for improvement in the reliability of its positive predictions.

6. Conclusions

In conclusion, as IoT continues to evolve and integrate deeper into our daily lives, ensuring the security of these devices becomes paramount. Through machine learning models and ensemble techniques, this notebook showcases a potential approach to enhancing IoT threat detection capabilities.